

# Dave530Week3

December 14, 2024

```
[111]: # 3.2 Exercise: Preparing for Exploratory Data Analysis Using Python
```

```
[112]: # Exercise 1: Page 11: 1-1
# Execute existing cells and solve exercise as directed.
```

```
[202]: import warnings
# Suppress all warnings
warnings.filterwarnings("ignore")
import nsfg
```

```
[204]: preg = nsfg.ReadFemPreg()
preg.head()
```

```
[204]:
```

	caseid	pregordr	howpreg_n	howpreg_p	moscurrp	nowprgdk	pregend1	\
0	1	1	NaN	NaN	NaN	NaN	6.0	
1	1	2	NaN	NaN	NaN	NaN	6.0	
2	2	1	NaN	NaN	NaN	NaN	5.0	
3	2	2	NaN	NaN	NaN	NaN	6.0	
4	2	3	NaN	NaN	NaN	NaN	6.0	

	pregend2	nbrnaliv	multbrth	...	laborfor_i	religion_i	metro_i	\
0	NaN	1.0	NaN	...	0	0	0	
1	NaN	1.0	NaN	...	0	0	0	
2	NaN	3.0	5.0	...	0	0	0	
3	NaN	1.0	NaN	...	0	0	0	
4	NaN	1.0	NaN	...	0	0	0	

	basewgt	adj_mod_basewgt	finalwgt	secu_p	sest	cmintvw	\
0	3410.389399	3869.349602	6448.271112	2	9	NaN	
1	3410.389399	3869.349602	6448.271112	2	9	NaN	
2	7226.301740	8567.549110	12999.542264	2	12	NaN	
3	7226.301740	8567.549110	12999.542264	2	12	NaN	
4	7226.301740	8567.549110	12999.542264	2	12	NaN	

	totalwgt_lb
0	8.8125
1	7.8750
2	9.1250

```
3      7.0000
4      6.1875
```

```
[5 rows x 244 columns]
```

Print the column names.

```
[205]: preg.columns
```

```
[205]: Index(['caseid', 'pregordr', 'howpreg_n', 'howpreg_p', 'moscurrp', 'nowprgdk',
          'pregend1', 'pregend2', 'nbrnaliv', 'multbrth',
          ...
          'laborfor_i', 'religion_i', 'metro_i', 'basewgt', 'adj_mod_basewgt',
          'finalwgt', 'secu_p', 'sest', 'cmintvw', 'totalwgt_lb'],
          dtype='object', length=244)
```

Select a single column name.

```
[206]: preg.columns[1]
```

```
[206]: 'pregordr'
```

Select a column and check what type it is.

```
[120]: pregordr = preg['pregordr']
       type(pregordr)
```

```
[120]: pandas.core.series.Series
```

Print a column.

```
[122]: pregordr
```

```
[122]: 0      1
      1      2
      2      1
      3      2
      4      3
      ..
    13588    1
    13589    2
    13590    3
    13591    4
    13592    5
      Name: pregordr, Length: 13593, dtype: int64
```

Select a single element from a column.

```
[124]: pregordr[0]
```

```
[124]: 1
```

Select a slice from a column.

```
[126]: pregordr[2:5]
```

```
[126]: 2    1
      3    2
      4    3
      Name: pregordr, dtype: int64
```

Select a column using dot notation.

```
[128]: pregordr = preg.pregordr
```

Count the number of times each value occurs.

```
[130]: preg.outcome.value_counts().sort_index()
```

```
[130]: outcome
      1    9148
      2    1862
      3     120
      4    1921
      5     190
      6     352
      Name: count, dtype: int64
```

Check the values of another variable.

```
[132]: preg.birthwgt_lb.value_counts().sort_index()
```

```
[132]: birthwgt_lb
      0.0      8
      1.0     40
      2.0     53
      3.0     98
      4.0    229
      5.0    697
      6.0   2223
      7.0   3049
      8.0   1889
      9.0    623
     10.0    132
     11.0     26
     12.0     10
     13.0      3
     14.0      3
     15.0      1
      Name: count, dtype: int64
```

Make a dictionary that maps from each respondent's `caseid` to a list of indices into the pregnancy `DataFrame`. Use it to select the pregnancy outcomes for a single respondent.

```
[134]: caseid = 10229
preg_map = nsfg.MakePregMap(preg)

#jyoti

indices = preg_map[caseid]
preg.outcome[indices].values
```

```
[134]: array([4, 4, 4, 4, 4, 4, 1], dtype=int64)
```

## 0.1 Exercises

Select the `birthord` column, print the value counts, and compare to results published in the [codebook](#)

```
[137]: birthord = preg['birthord']
birthord
```

```
[137]: 0      1.0
      1      2.0
      2      1.0
      3      2.0
      4      3.0
      ...
      13588  1.0
      13589  NaN
      13590  NaN
      13591  2.0
      13592  3.0
      Name: birthord, Length: 13593, dtype: float64
```

We can also use `isnull` to count the number of nans.

```
[139]: preg.birthord.isnull().sum()
```

```
[139]: 4445
```

Select the `prglngth` column, print the value counts, and compare to results published in the [codebook](#)

```
[141]: prglngth = preg['prglngth']
print(prglngth)
preg.prglngth.value_counts().sort_index()
```

```
0      39
1      39
2      39
```

```

3      39
4      39
..
13588  39
13589   6
13590   5
13591  39
13592  39
Name: prglength, Length: 13593, dtype: int64

```

```

[141]: prglength
0      15
1       9
2      78
3     151
4     412
5     181
6     543
7     175
8     409
9     594
10     137
11     202
12     170
13     446
14      29
15      39
16      44
17     253
18      17
19      34
20      18
21      37
22     147
23      12
24      31
25      15
26     117
27       8
28      38
29      23
30     198
31      29
32     122
33      50
34      60
35     357

```

```

36      329
37      457
38      609
39     4744
40     1120
41      591
42      328
43      148
44       46
45       10
46        1
47        1
48        7
50        2
Name: count, dtype: int64

```

To compute the mean of a column, you can invoke the `mean` method on a Series. For example, here is the mean birthweight in pounds:

```
[143]: preg.totalwgt_lb.mean()
```

```
[143]: 7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
[145]: # homework
# convert totalwgt_lb to totalwgt_kg and create a new column
preg['totalwgt_kg'] = preg.totalwgt_lb * 0.453592
preg.totalwgt_kg.mean()
```

```
[145]: 3.2956309433502984
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
[147]: resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
[149]: resp.head()
```

```
[149]:
```

	caseid	rscrinf	rdormres	rostscrn	rscreenhisp	rscreenrace	age_a	\
0	2298	1	5	5	1	5.0	27	
1	5012	1	5	1	5	5.0	42	
2	11586	1	5	1	5	5.0	43	
3	6794	5	5	4	1	5.0	15	
4	616	1	5	4	1	5.0	20	

	age_r	cmbirth	agescrn	...	pubassis_i	basewgt	adj_mod_basewgt	\
0	27	902	27	...	0	3247.916977	5123.759559	
1	42	718	42	...	0	2335.279149	2846.799490	
2	43	708	43	...	0	2335.279149	2846.799490	
3	15	1042	15	...	0	3783.152221	5071.464231	
4	20	991	20	...	0	5341.329968	6437.335772	

	finalwgt	secu_r	sest	cmintvw	cmlstyr	screentime	intvlngh
0	5556.717241	2	18	1234	1222	18:26:36	110.492667
1	4744.191350	2	18	1233	1221	16:30:59	64.294000
2	4744.191350	2	18	1234	1222	18:19:09	75.149167
3	5923.977368	2	18	1234	1222	15:54:43	28.642833
4	7229.128072	2	18	1233	1221	14:19:44	69.502667

[5 rows x 3087 columns]

Select the `age_r` column from `resp` and print the value counts. How old are the youngest and oldest respondents?

```
[151]: # homework
# Select the age_r column from resp and print the value counts
resp.age_r.value_counts().sort_index()
```

```
[151]: age_r
15    217
16    223
17    234
18    235
19    241
20    258
21    267
22    287
23    282
24    269
25    267
26    260
27    255
28    252
29    262
30    292
31    278
32    273
33    257
34    255
35    262
36    266
37    271
```

```

38    256
39    215
40    256
41    250
42    215
43    253
44    235
Name: count, dtype: int64

```

```
[152]: # From the above result, The youngest one is 15 year old and oldest one is 44
```

We can use the `caseid` to match up rows from `resp` and `preg`. For example, we can select the row from `resp` for `caseid` 2298 like this:

```
[154]: resp[resp.caseid==2298]
```

```

[154]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298         1         5         5           1         5.0    27

   age_r  cmbirth  agescrn  ...  pubassis_i      basewgt  adj_mod_basewgt  \
0     27     902      27  ...           0  3247.916977    5123.759559

   finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngh
0  5556.717241      2   18    1234    1222    18:26:36  110.492667

[1 rows x 3087 columns]

```

And we can get the corresponding rows from `preg` like this:

```
[156]: preg[preg.caseid==2298]
```

```

[156]:   caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk  pregend1  \
2610    2298         1        NaN        NaN        NaN        NaN        6.0
2611    2298         2        NaN        NaN        NaN        NaN        6.0
2612    2298         3        NaN        NaN        NaN        NaN        6.0
2613    2298         4        NaN        NaN        NaN        NaN        6.0

   pregend2  nbrnaliv  multbrth  ...  religion_i  metro_i      basewgt  \
2610      NaN        1.0      NaN  ...         0         0  3247.916977
2611      NaN        1.0      NaN  ...         0         0  3247.916977
2612      NaN        1.0      NaN  ...         0         0  3247.916977
2613      NaN        1.0      NaN  ...         0         0  3247.916977

   adj_mod_basewgt  finalwgt  secu_p  sest  cmintvw  totalwgt_lb  \
2610    5123.759559  5556.717241      2   18      NaN        6.8750
2611    5123.759559  5556.717241      2   18      NaN        5.5000
2612    5123.759559  5556.717241      2   18      NaN        4.1875
2613    5123.759559  5556.717241      2   18      NaN        6.8750

```



	totalwgt_kg
2610	3.118445
2611	2.494756
2612	1.899417
2613	3.118445

[4 rows x 245 columns]

How old is the respondent with caseid 1?

```
[158]: #homework
# filter the age based on caseid == 1

resp[resp.caseid==1].age_r
```

```
[158]: 1069    44
Name: age_r, dtype: int64
```

What are the pregnancy lengths for the respondent with caseid 2298?

```
[160]: #homework
# filter data based on caseid

preg[preg.caseid==2298].prglngth
```

```
[160]: 2610    40
2611    36
2612    30
2613    40
Name: prglngth, dtype: int64
```

What was the birthweight of the first baby born to the respondent with caseid 5013?

```
[162]: # homework

# use `pregordr` to select the first baby

preg[(preg.caseid==5013) & (preg.pregordr==1)].birthwgt_lb
```

```
[162]: 5516    7.0
Name: birthwgt_lb, dtype: float64
```

```
[163]: # # Exercise 2: Page 11: 1-2
# Find the value of pregnum
```

```
[164]: pregnum = preg.groupby('caseid').size()

# Display the first few rows
print(pregnum)
```

```

# Convert the result to a DataFrame
pregnancy_count_df = pregnum.reset_index(name='pregnum')

# Count how many respondents have the same pregnancy count
same_pregnancy_count = pregnancy_count_df['pregnum'].value_counts().
↳reset_index()
same_pregnancy_count.columns = ['pregnum', 'respondent_count']

# Sort by pregnancy_count
same_pregnancy_count = same_pregnancy_count.sort_values(by='pregnum')
# Display the result
print(same_pregnancy_count)

```

```

caseid
1      2
2      3
6      3
7      2
12     1
..
12566  2
12568  3
12569  2
12570  3
12571  5
Length: 5033, dtype: int64
   pregnum  respondent_count
1         1             1267
0         2             1432
2         3             1110
3         4              611
4         5              305
5         6              150
6         7              80
7         8              40
8         9              21
9        10               9
10       11               3
11       12               2
12       14               2
13       19               1

```

[165]: # To compare groups using summary statistics in Python, you can calculate  
↳descriptive statistics like  
#mean, median, standard deviation, count, etc., for each group. Using the  
↳pandas library, this can be done efficiently.

```
[166]: import numpy as np
      # import math as math
      import thinkstats2
      import matplotlib.pyplot as plt
      import numpy as np
      import pandas as pd
      %matplotlib inline
      import thinkplot
```

```
[167]: # # Exercise 3: Page 25: 2-1
      # Do first baby arrive late?
```

```
[168]: import nsfg
```


```
[169]: preg = nsfg.ReadFemPreg()
      live = preg[preg.outcome == 1]
```

```
[170]: ages = np.floor(live.agepreg)
```

```
[171]: # As an exercise, plot the histogram of pregnancy lengths (column prglngth).
```

```
[172]: # Solution


hist = thinkstats2.Hist(live.prglngth, label='prglngth')
thinkplot.Hist(hist)
thinkplot.Config(xlabel='weeks', ylabel='Count')
```



```
[173]: # From live births, we can select first babies and others using birthord,  
  
# List item  
# List item  
# then compute histograms of pregnancy length for the two groups.
```

```
[174]: firsts = live[live.birthord == 1]  
others = live[live.birthord != 1]  
  
first_hist = thinkstats2.Hist(firsts.prglnth, label='first')  
other_hist = thinkstats2.Hist(others.prglnth, label='other')
```

```
[175]: width = 0.45  
thinkplot.PrePlot(2)  
thinkplot.Hist(first_hist, align='right', width=width)  
thinkplot.Hist(other_hist, align='left', width=width)  
thinkplot.Config(xlabel='weeks', ylabel='Count', xlim=[27, 46])
```



```
[176]: #Series provides methods to compute summary statistics:
```

```
mean = live.prglnth.mean()
var = live.prglnth.var()
std = live.prglnth.std()
```

```
[177]: mean, std
```

```
[177]: (38.56055968517709, 2.702343810070593)
```

```
[178]: # Here's are the mean pregnancy lengths for first babies and others:
```

```
# List item
# List item
```

```
[179]: firsts.prglnth.mean(), others.prglnth.mean()
```

```
[179]: (38.60095173351461, 38.52291446673706)
```

```
[180]: #Here's the difference (in weeks):
```

```
firsts.prglnth.mean() - others.prglnth.mean()
```

[180]: 0.07803726677754952

```
[181]: # Distributions
# Summary statistics - mean, variance, median, etc. Dangerous because they
#       obscure the data.
# Alternative is to look at distributions of the data - describes how often
#       each value appears.

# Distributions are usually represented as histograms (raw frequencies binned
#       into equally spaced buckets).
# Normalized histogram is called Probability Mass Function (PMF).
```

```
[182]: prglen_first_babies = firsts.prglength
prglen_other_babies = others.prglength

# setting up range of histogram and number of bins
first_baby_min_prglen = np.min(prglen_first_babies)
first_baby_max_prglen = np.max(prglen_first_babies)
other_baby_min_prglen = np.min(prglen_other_babies)
other_baby_max_prglen = np.max(prglen_other_babies)
print("first baby preg length min: %d, max: %d" %
      (first_baby_min_prglen, first_baby_max_prglen))
print("other baby preg length min: %d, max: %d" %
      (other_baby_min_prglen, other_baby_max_prglen))

bin_lb = min([first_baby_min_prglen, other_baby_min_prglen])
bin_ub = max([first_baby_max_prglen, other_baby_max_prglen])
nbr_bins = bin_ub - bin_lb
bin_range = (bin_lb, bin_ub)
print("range:", bin_range, "#-bins:", nbr_bins)

# building the histograms
first_baby_fdist = np.histogram(np.array(prglen_first_babies), bins=nbr_bins,
#       range=bin_range)
other_baby_fdist = np.histogram(np.array(prglen_other_babies), bins=nbr_bins,
#       range=bin_range)
```

```
first baby preg length min: 0, max: 48
other baby preg length min: 4, max: 50
range: (0, 50) #-bins: 50
```

```
[183]: def mode(fdist):
        """ takes a histogram and returns the most frequent value """
        mode_idx = np.argmax(fdist[0])
        return fdist[1][mode_idx]

def all_modes(fdist):
```

```

""" takes a histogram and returns (value, freq) pairs in desc freq """
mode_idx = np.argsort(fdist[0])[:, :-1]
vf_pairs = []
for i in range(mode_idx.shape[0]):
    vf_pairs.append((int(fdist[1][mode_idx[i]]), fdist[0][mode_idx[i]]))
return vf_pairs

```

```

[184]: print("First baby arrival top week (mode): %d" % (mode(first_baby_fdist)))
print("Other baby arrival top week (mode): %d" % (mode(other_baby_fdist)))

print("First baby top 5 frequent weeks:", all_modes(first_baby_fdist)[0:5])
print("Other baby top 5 frequent weeks:", all_modes(other_baby_fdist)[0:5])

```

```

First baby arrival top week (mode): 39
Other baby arrival top week (mode): 39
First baby top 5 frequent weeks: [(39, 2114), (40, 536), (41, 360), (38, 272),
(37, 208)]
Other baby top 5 frequent weeks: [(39, 2579), (40, 580), (38, 335), (37, 247),
(41, 227)]

```

```

[185]: # Plotting Histograms
# From histogram below, it looks like until about week 41, other babies tend to
    ↪ arrive sooner in aggregate,
#but after that more first babies arrive. So there seems to be evidence that
    ↪ first babies do tend to arrive late.

# Other information from the plot:

# Mode - most common week for arrival seems to be 39 weeks.
# Shape - distribution is assymetric around the mode, drops of quicker to right
    ↪ than left.
#Possible reason is because surgical intervention happens for late pregnancies.
# Outliers - babies born at 30 weeks are probably just unusual, but may also
    ↪ reflect reporting errors.

```

```

[186]: def remaining_lifetime(pmf, age):
    """ Takes PMF and age and returns PMF of remaining lifetime """
    ages = pmf[1][: -1]
    idxs = np.where(ages >= age)[0]
    r_probs = pmf[0][idxs]
    r_ages = ages[idxs]
    # generate a fake population with r_probs and r_ages
    fake_pop_size = 1000
    fake_pop = []
    for i in range(r_ages.shape[0]):
        num_aged = int(r_probs[i] * fake_pop_size)
        for j in range(num_aged):

```

```

        fake_pop.append(r_ages[i])
    # create a histogram and return
    r_nbins = r_ages.shape[0]
    r_range = (np.min(r_ages), np.max(r_ages))
    return np.histogram(np.array(fake_pop), bins=r_nbins, range=r_range)

```

```

[187]: # compute PMFs for each distribution
pmf_first_babies = np.histogram(prglen_first_babies,
                                bins=nr_bins, range=bin_range)
pmf_other_babies = np.histogram(prglen_other_babies,
                                bins=nr_bins, range=bin_range)

def cond_prob(pmf, x):
    """ Returns probability that birth will occur at week x given person
        is pregnant at week x """
    return remaining_lifetime(pmf, x)[0][0]


first_birth_probs = []
other_birth_probs = []
for x in range(0, 49):
    first_birth_probs.append(cond_prob(pmf_first_babies, x))
    other_birth_probs.append(cond_prob(pmf_other_babies, x))

plt.plot(range(len(first_birth_probs)), first_birth_probs, color='r',
         label="first births")
plt.plot(range(len(first_birth_probs)), other_birth_probs, color='b',
         label="other births")

plt.ylabel("probability(birth|pregnant for x weeks)")
plt.xlabel("weeks")
plt.xlim([25, 50])
plt.legend(loc="best")
plt.grid()
plt.show()

```





[188]: # The chart indicates that a first baby is less likely to be born until about week 42, after which they are more likely to be born than other babies.

[189]: # Reporting Results  
# The evidence clearly indicates that first babies tend to arrive later than others.

[190]: # Exercise 4: Page 25: 2-4  
# Using the variable totalwgt\_lb, investigate whether first babies are lighter or heavier than others.  
  
# Compute Cohen's effect size to quantify the difference between the groups. How does it compare to the difference in pregnancy length?

[191]: # Compute the Cohen effect size for the difference in pregnancy length for first babies and others.  
# The Cohen's Effect Size measures the magnitude of the difference between two groups.

```
[192]: #This function computes the Cohen effect size, which is the difference in means  
↪expressed in number of standard deviations:
```

```
[193]: def CohenEffectSize(group1, group2):  
    """Computes Cohen's effect size for two groups.  
  
group1: Series or DataFrame  
group2: Series or DataFrame  
  
returns: float if the arguments are Series;  
Series if the arguments are DataFrames  
    """  
    diff = group1.mean() - group2.mean()  
  
    var1 = group1.var()  
    var2 = group2.var()  
    n1, n2 = len(group1), len(group2)  
  
    pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)  
    d = diff / np.sqrt(pooled_var)  
    return d
```

```
[194]: # mean1, mean2: Compute the means of the two groups.  
# std1, std2: Compute the standard deviations of the two groups (with ddof=1  
↪for sample standard deviation).  
# pooled_std: Calculate the pooled standard deviation using the formula.  
# d: Compute Cohen's  
# d by dividing the difference in means by the pooled standard deviation.  
  
# Interpretation of Cohen's  
# d=0.2: Small effect size  
# d=0.5: Medium effect size  
# d=0.8: Large effect size
```

```
[195]: # Solution  
# Find the mean of total weight of the first and other babies in lb  
  
firsts.totalwgt_lb.mean(), others.totalwgt_lb.mean()
```

```
[195]: (7.201094430437772, 7.325855614973262)
```

```
[196]: # Find the Cohen'd effect size  
  
CohenEffectSize(firsts.totalwgt_lb, others.totalwgt_lb)
```

```
[196]: -0.088672927072602
```

```
[197]: # Conclusion: A Cohen's d of
# -0.088672927072602 suggests a very small difference between the firsts and
# other babies group, with the other babies group being slightly higher.
# A negative indicates that the mean of the other babies total wt group is
# higher than the mean of the first total wt group.
```

```
[198]: # How does it compare to the difference in pregnancy length?
```

```
[199]: # Solution
firsts.prglength.mean(), others.prglength.mean()
```

```
[199]: (38.60095173351461, 38.52291446673706)
```

```
[200]: CohenEffectSize(firsts.prglength, others.prglength)
```

```
[200]: 0.028879044654449883
```

```
[201]: # Conclusion: A Cohen's d of
# d=0.028879044654449883 falls well within the small effect size range and is
# very close to zero, it indicates almost no meaningful difference between the
# two groups.
# A positive d means the mean of prglength of first group is slightly higher
# than the mean of prglength of others group.
```